



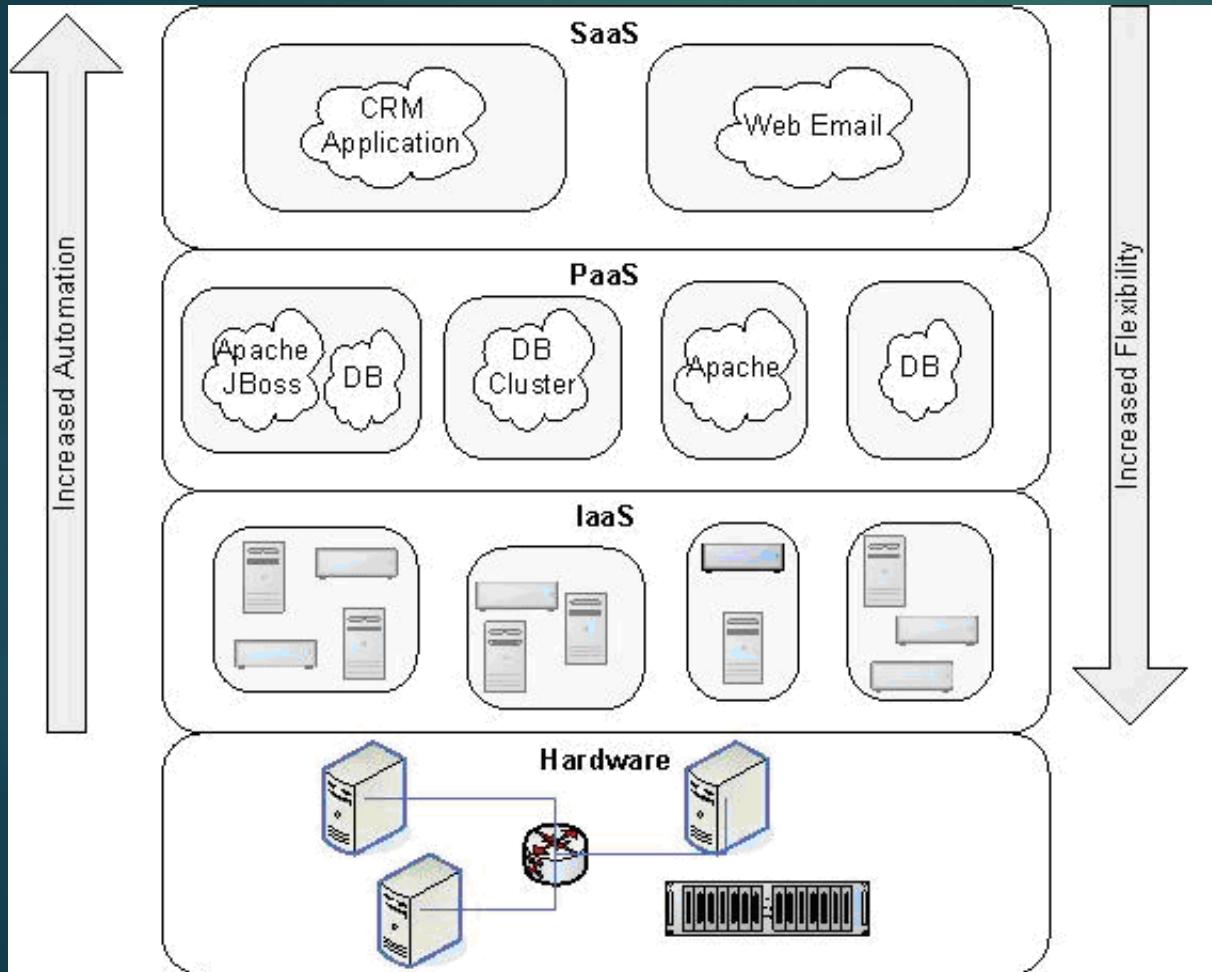
Software As A Service and Multitenancy

K V SUBRAMANIAM



Background

Comparison: Cloud Service Models

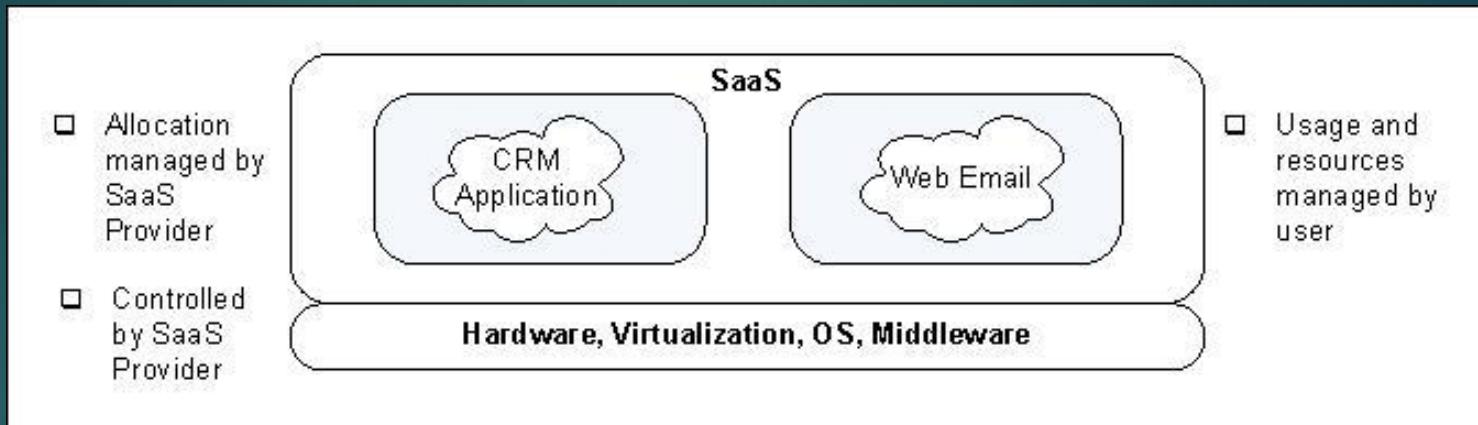


Software As A Service (SaaS)

Platform As A Service (PaaS)

Infrastructure As A Service (IaaS)

Software as a Service (SaaS)



NIST definition: The capability provided to the consumer is to use the provider's applications running on a cloud infrastructure. The applications are accessible from various client devices through a thin client interface such as a web browser.

SaaS Benefits

Characteristics	Benefits
Network delivered access to commercially available software	No local infrastructure or software to purchase or maintain Applications & data are available anywhere with network connectivity
Application delivery is one-to-many model	Operating costs are reduced by managing infrastructure in central locations rather than at each customer's site
Built on optimized & robust platform	Improved availability and reliability
Customer pays for as much as they need when they need it	Lower TCO

Exercise 0: Revenue Options

- ▶ Subscription (monthly rental)
 - ▶ Freemium model – starts free and then increases price
 - ▶ Tiered pricing
 - ▶ Usage/transaction based
 - ▶ Per user pricing
 - ▶ Per feature pricing
 - ▶ Ad-based revenue (pay per click)
 - ▶ Profit sharing
 - ▶ Ownership Sharing
- ▶ For each of the options on the left, name one SaaS company that does this

SaaS architecture

- ▶ Front-end or GUI for SaaS application is a browser
- ▶ Business Logic runs on the backend cloud infrastructure
- ▶ Business Logic frequently implemented as *microservices*
- ▶ Examples: Google Docs, Gmail, Salesforce.com

Exercise 1

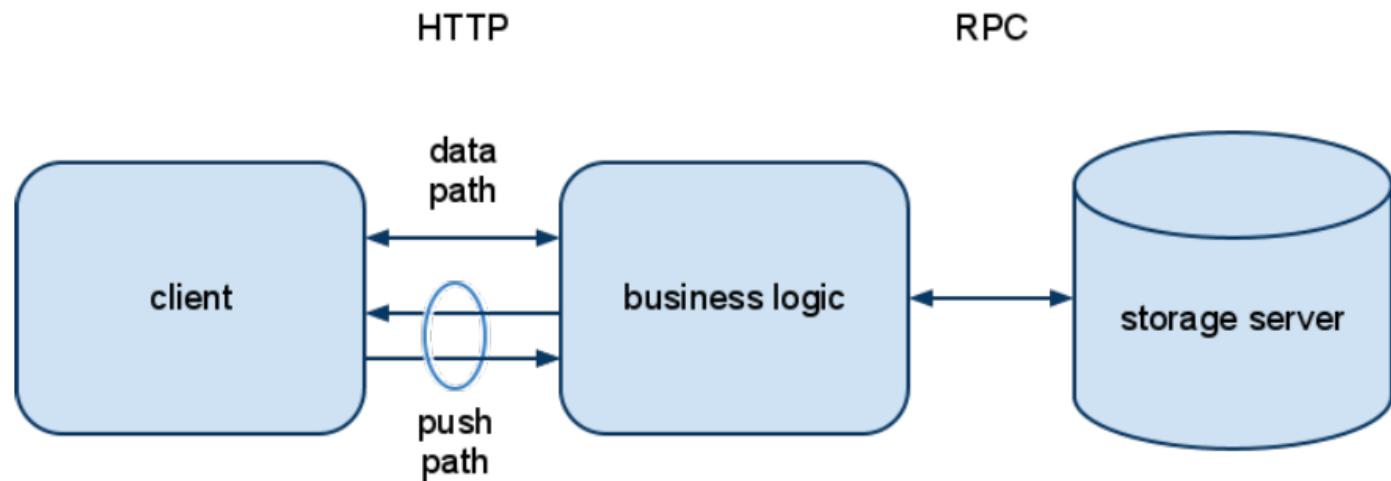
- ▶ You have built a Bookstore application - which permits users to login to a portal to query and purchase books
- ▶ What UI challenges will you have to face?

Solution GUI

- ▶ Need screens similar to Amazon for each microservice
- ▶ Challenges
 - ▶ Ensure that the app runs well on all browsers
 - ▶ Handling performance lags

Case Study: Gmail

Macro-architecture



data path: code, styles, preferences, user data, ...

push path: chat, new mail, presence, contact updates, buzz updates, ...



Gmail: Client (Front-end) Functions

- ▶ Look and feel
 - ▶ Builds all UI based on browser type
 - ▶ Loads GUI code (Javascript) as needed
 - ▶ Widgets: Provides uniform look/feel to different windows
- ▶ Performance
 - ▶ Fetches and caches data
 - ▶ Records actions for GUI performance analysis
 - ▶ Determine round-trip time for a 1pixel image
- ▶ Reports presence/idle to back end
- ▶ Drives multiple windows

Exercise 2 (10 mins)

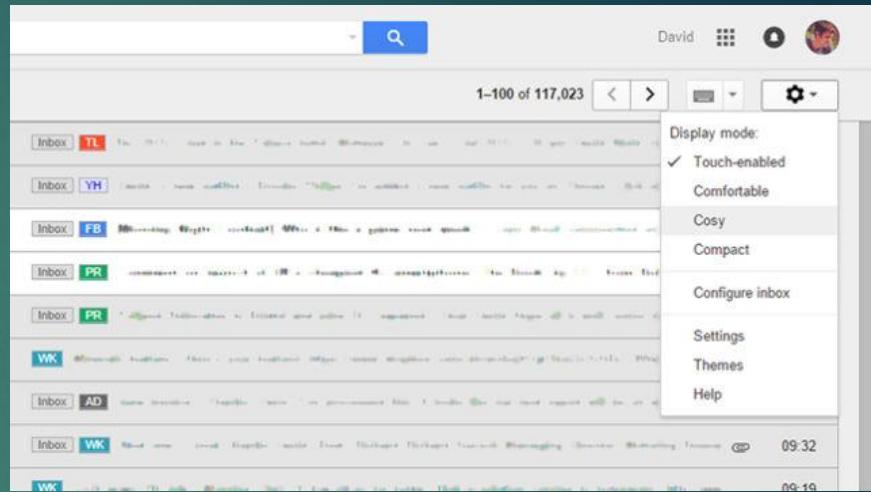
- ▶ For the bookkarts application which permits users to login to a portal to query and purchase books
- ▶ Pustakhouse buys this application
- ▶ Approached by another bookstore – ClassyBooks that wants additionally
 - ▶ Name + Front cover + review of books to be displayed
 - ▶ Different layout of the page
 - ▶ Customers to be able to make one click payment
- ▶ How will you customize the GUI?

Solution: GUI Screens

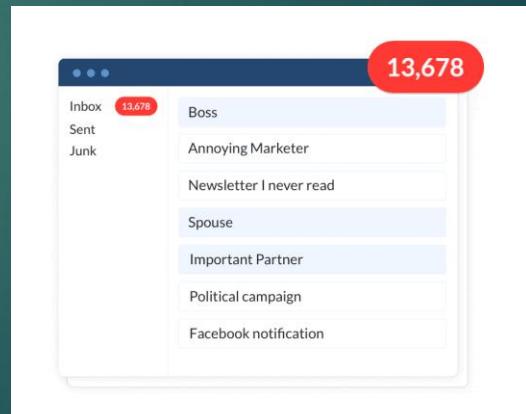
- ▶ How to customize the screen?
 - ▶ Typically we need one GUI for the app as used by the customers of the SaaS application
 - ▶ The business application
 - ▶ Another one for customizing the App for the vendor for use by the vendor
 - ▶ Used for defining the layout, logos

Case Study: Gmail Customization

- ▶ Settings and Layout
- ▶ Gmail Themes
- ▶ Third party add ons
 - ▶ Sane-Box



Further Reading – Details of Gmail architecture



Exercise 3 (10 mins)

- ▶ For the bookkarts application which permits users to login to a portal to query and purchase books
- ▶ Pustakhouse buys this application
- ▶ Approached by another bookstore – ClassyBooks that wants additionally
 - ▶ Name + Front cover + review of books to be displayed
 - ▶ Different layout of the page
 - ▶ Customers to be able to make one click payment
 - ▶ Different schemas?
- ▶ How will you design the backend?

Exercise 3 – Backend Design

- ▶ Name + Front cover + review of books to be display
- ▶ Customers to be able to make one click payment
 - ▶ Requires different workflows
 - ▶ Allow users to customize workflows
 - ▶ Separate functionality into microservices – RESTful services
 - ▶ Different layout of the page
 - ▶ Allow GUI customization
 - ▶ Different schemas
 - ▶ Multitenancy.

See details about Backend Design
in the Further Reading section

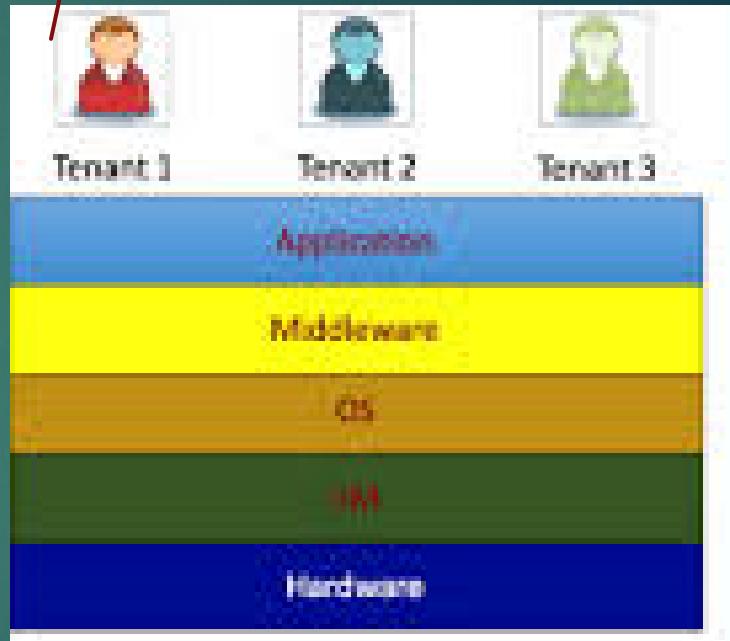
Multi-Tenancy

Multi-tenancy vs Virtualization

- ▶ Virtualization: coarse-grained resource sharing
 - ▶ Generally hardware, difficult to do fine-grained sharing
 - ▶ Compute / cores
 - ▶ Memory
- ▶ Fine-grained resource sharing, generally software

Terminology

- ▶ Tenants
 - ▶ Customers of the cloud provider
 - ▶ Organizations or users
- ▶ Users
 - ▶ Users of the cloud system
 - ▶ May belong to a tenant



Think of tenant as a family that rents an apartment and its services like gym etc.

Users are people that stay in an apartment.

Motivation

- ▶ Experiment by Jacobs et al.
- ▶ 3 methods for creating a multi-tenant database
 - ▶ Shared machine
 - ▶ 55 MB main memory, 4MB disk memory per tenant
 - ▶ Shared process
 - ▶ 80 MB main memory for 10,000 tenants
 - ▶ Challenges
 - ▶ Security
 - ▶ Customization
 - ▶ Different schema
- ▶ Shared machine
 - ▶ Db process and tables per tenant
- ▶ Shared process
 - ▶ DB tables per tenant
- ▶ Shared table
 - ▶ Process and tables shared

Jacobs, D, Aulbach, S, Marz 2007. Ruminations on Multi-Tenant Databases. 12.GI-Fachtagung fur Datenbanksysteme in Business, Technologie und Web (BTW 2007), 5 bis 9, Aachen, Germany.

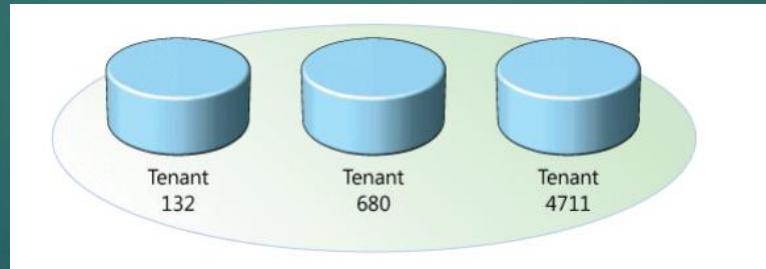
Multi-tenancy Levels - Summary

- ▶ Different types of multitenancy possible
- ▶ Simplest - each client has own version of software
- ▶ Most complex - single version of software hosted on a cluster.
Supports multiple tenants.

Further Reading section – for details

Storage Multi-tenancy

- ▶ File systems already have well-known methods for sharing
- ▶ Database methods
 - ▶ Dedicated tables per client
 - ▶ Shared tables per client



Dedicated Tables per Tenant

- ▶ Each tenant has their own tables
- ▶ Shared table, separate schema
- ▶ Generally DBs store different tables in different files
 - ▶ Possibly greater security
- ▶ Details
 - ▶ Register the three garages as DB users
 - ▶ GRANT SELECT ON FriendlyTable TO FriendlyGarage WITH GRANT OPTION

Best Garage		
Car License	Service	Cost

Friendly Garage		
Car License	Service	Cost

Honest Garage		
Car License	Service	Cost

Shared Table Option - Basics

- ▶ All data stored in the same table
- ▶ Each row has a TenantId which identifies the tenant to which the data belongs
- ▶ More space efficient, but may need an additional SELECT
- ▶ Can use a key to encrypt the data for each tenant
- ▶ Metadata table to store information about the client

Data Table 1			
Tenant Id	Car License	Repair	Cost
1			
2			
2			
1			
3			
2			

Metadata Table 1	
Tenant Id	Data
1	Best Garage
2	Friendly Garage
3	Honest Garage

Shared Table Option - Customization

- ▶ Customization
 - ▶ Different tenants would want to define their own custom fields
 - ▶ How can this be done efficiently?

DB Customization – Pre-allocated Columns

- ▶ A number of pre-defined custom columns in each row
 - ▶ Salesforce has 500
- ▶ Metadata table stores the type and name of each column

Data Table 1

Tenant Id	Car License	Service	Cost	Custom1	Custom2
1					
2					
2					
1					
3					
2					

Metadata Table 1

Tenant Id	Tenant Name	Custom1 Name	Custom1 Type
1	Best Garage	Service Rating	int
2	Friendly Garage	Service Manager	string
3	Honest Garage		

Exercise 4 (5 minutes)

- ▶ Consider the table to the right
- ▶ Consider a query from *BestGarage*
 - ▶ *SELECT CarLicense, ServiceRating FROM CustomerTable WHERE ServiceRating > 8;*
- ▶ What
 - ▶ Algorithm would be used by the compiler to translate this into a query on *DataTable1*?
 - ▶ What is the resulting query?

Data Table 1

Tenant Id	Car License	Service	Cost	Custom1	Custom2
1					
2					
2					
1					
3					
2					

Metadata Table 1

Tenant Id	Tenant Name	Custom1 Name	Custom1 Type
1	Best Garage	Service Rating	int
2	Friendly Garage	Service Manager	string
3	Honest Garage		

Exercise 4 (5 minutes)

- ▶ Consider query
 - ▶ `SELECT CarLicense,
ServiceRating FROM
CustomerTable WHERE
ServiceRating > 8;`
- ▶ Algorithm
 - ▶ Check from *Metadata Table* that *ServiceRating* field exists for *BestGarage*
 - ▶ Find the corresponding column from *MetadataTable 1* as *Column 1*
 - ▶ Add an additional condition to *WHERE* clause
 - ▶ `SELECT CarLicense,
int(Custom1) FROM
CustomerTable WHERE
(int(Custom1) > 8) AND
(TenantId=1);`

Data Table 1

Tenant Id	Car License	Service	Cost	Custom1	Custom2
1					
2					
2					
1					
3					
2					

Metadata Table 1

Tenant Id	Tenant Name	Custom1 Name	Custom1 Type
1	Best Garage	Service Rating	int
2	Friendly Garage	Service Manager	string
3	Honest Garage		

DB Customization – Name-value Pairs

- ▶ Main table has a name-value pair record column
- ▶ Pivot table has list of name value pair
 - ▶ If there is more than one name-value pair, there will be multiple records in the pivot table
- ▶ Name is actually an id that points to the actual name and type in a metadata table
- ▶ Greater space efficiency; however joins may be needed to re-construct records

The diagram illustrates the schema for managing name-value pairs. It consists of four tables: Data Table 1, Data Table 2 (Name Value Pairs), Metadata Table 1, and Metadata Table 2.

Data Table 1:

Tenant Id	Car License	Service	Cost	NameValuePairRec
1				275
2				
2				
1				
3				
2				

Data Table 2 (Name Value Pairs):

NameValuePairRec	NameID	Value
275	15	5.5

Metadata Table 1:

Name Id	Name	Type
15	Service Rating	int
	Service Manager	string

Metadata Table 2:

Tenant Id	Data
1	Best Garage
2	Friendly Garage
3	Honest Garage

Exercise 5 (15 minutes)

- ▶ Suppose *BestGarage* has an additional custom field *MechanicId* of type integer
 - ▶ How would *DataTable2* and *MetadataTable1* change?
 - ▶ Repeat previous exercise for *BestGarage* query
 - ▶ *SELECT CarLicense, MechanicId, ServiceRating FROM DataTable WHERE ServiceRating > 8*

Data Table 1

Tenant Id	Car License	Service	Cost	NameValuePairRec
1				275
2				
2				
1				
3				
2				

Data Table 2 (Name Value Pairs)

NameValuePairRec	NameID	Value
275	15	5.5

Metadata Table 1

Name Id	Name	Type
15	Service Rating	int
	Service Manager	string

Metadata Table 2

Tenant Id	Data
1	Best Garage
2	Friendly Garage
3	Honest Garage

Scaling in SaaS

- ▶ Built on PaaS
- ▶ So, all PaaS based scalability is relevant to SaaS
- ▶ Additionally
 - ▶ Multi-tenancy for improved scaling/utilization
 - ▶ Query optimization – per tenant in real time
 - ▶ <300 ms latency

SaaS Summary

- ▶ SaaS based on the following principles
 - ▶ Multi-tenancy
 - ▶ Customization
 - ▶ Scaling