



Partitioning in Cloud Storage

K V SUBRAMANIAM

CHAPTER 5 : REPLICATION – MARTIN KLEPPMAN – DESIGNING DATA INTENSIVE APPLICATIONS

What is a partition


- ▶ Definition
- ▶ “Mathematically, a partition refers to dividing a set or a number into disjoint, non-empty subsets that together cover the entire original set or sum to the original number.”
- ▶ When your dataset is broken into disjoint non-empty subsets.
- ▶ Each subset can be independently managed

Why partition

- Cloud applications scale using elastic resources,
 - backend bottlenecks can occur if the data store remains the same.
- Partitioning
 - divides large data into smaller parts distributed across nodes,
 - improving query and IO performance.
- Each data item belongs to one partition,
 - allowing simultaneous operations on multiple partitions.
- Nodes process their own partitions independently,
 - enabling throughput scaling by adding nodes.
- For example, spreading data across multiple disks enhances overall IO performance.

Leader based replication

- Large, complex queries or heavy IO tasks can be parallelized across multiple nodes.
- Partitioning is often paired with replication, so each partition's copies exist on several nodes for fault tolerance.
- A single node can hold multiple partitions.



Goal of partitioning

- **Objective :**
 - spread data across nodes
 - so that query load is evenly distributed
- Fair share – total load/#partitions
 - Ideal case when data is equally distributed
 - Results in linear scaling with #partitions
- **Skew:** when some partition has more data than others
 - Increased skew reduces effect of partitioning
 - Extreme case – single node has all the data and load
- **Hot Spot**
 - Partition with disproportionately high load

Class exercise

- Consider a scenario where we have 1000 records and 4 nodes. Each node should be responsible for $\frac{1}{4}$ of the records.
- Design a scheme to partition the data into the 4 nodes

Partitioning schemes

- Key Range Partition
- Hash Partition

Key Range Partition

- **How**
 - Non-overlapping ranges of key to define subsets
 - Partition to node mapping for query handling
- Benefit
 - Easy for clients to process
- Prone to skew
 - Certain ranges can be more frequent
- Used by Hbase
- Can cause hotspots if key not chosen carefully – example: timestamps




Figure 6-2. A print encyclopedia is partitioned by key range.

Hash based partition

- **How**
 - HashFn(Key)
 - Assign range of hashes to a partition.
 - Good hash function distributes keys evenly
- Benefit
 - Solves the skew problem
- Disadvantages
 - Range queries are more difficult to implement




Figure 6-3. Partitioning by hash of key.

Rebalancing Partitions

- Due to changes over time
 - Query load increases
 - Data set size increases
 - Failures to nodes
- Requires data to be moved from one node to another
 - Requirements
 - Minimal data transfer
 - Read/Writes should continue while rebalancing is in progress
 - After rebalancing, skew is minimized

Rebalancing strategies

- ▶ What to avoid?
 - ▶ Hash Mod n
 - ▶ Rather assign ranges
 - ▶ Otherwise there will be movements across all nodes
- ▶ What will work?
 - ▶ Fixed number of partitions
 - ▶ Create more partitions than nodes
 - ▶ Map multiple partitions to a node.
 - ▶ Entire partition is moved





Figure 6-6. Adding a new node to a database cluster with multiple partitions per node.

Reading your own writes consistency

- Read-after-write consistency (read-your-writes consistency)
 - ensures that users always see their own updates when they read data again.
- It does not guarantee
 - updates from other users will be immediately visible.




Other partitioning strategies

- **Dynamic partitioning**
 - When partition grow beyond a threshold
 - Split it in two.
 - Can also be merged in case of deletes
 - Each node can handle multiple partitions.


Ring Consistent Hashing

Basics



- Hash disk id to find position of disk in ring
- Object stored in next disk in ring
- Disk id could be IP address, drive name, or combination

Adding a New Drive



- Objects in purple hash range have to be moved from disk 4 to 5
- Move from only 1 disk
- With large number (say 100) of disks, disks will be uniformly distributed

Consistent Hashing: Mapping between range of hash and a Drive can be variable.
Locating drive for a object requires hashing and then looking for the next drive clockwise

Modified Consistent Hashing


- ▶ Problem: May take some time to move a hash range
 - ▶ Objects may need to be marked unavailable when hash range is in transit
 - ▶ Locating an object also takes time.
- ▶ Solution
 - ▶ Divide hash range into *partitions* (*fixed size*)
 - ▶ These are not disk partitions
 - ▶ Move partitions from one disk to another

MCH: Partition Power

- ▶ Total partitions in cluster = $2^{\text{partition power}}$
- ▶ If partition power = 15 ,
 - ▶ Total partitions $2^{15} = 32,768$.
- ▶ Those 32,768 partitions mapped to the drives.
- ▶ Number of drives might change
- ▶ Number of partitions is fixed
 - ▶ Hash(object) → partition
 - ▶ Partition number does not change with increasing #drives
 - ▶ Only some partitions are moved to the new drive

Swift Partitions - 1 Node

Node 1



8 Disks - 16 Partitions/Disk

Example:


Assuming equally weighted disks.

$$8 * 16 = 128 \text{ partitions}$$

Property of SwiftStack Inc.


Swift Partitions - Adding A Node: Partitions Are Reassigned

Node 1



8 Disks - 8 Partitions/Disk

Node 2




+ 8 Disks - 8 Partitions/Disk

$$16 * 8 = 128 \text{ partitions}$$

Property of SwiftStack Inc.

MCH: Replication

- ▶ Partitions are replicated, not files
- ▶ Replica count determines number of replicas
- ▶ *Unique-as-possible* algorithm used
 - ▶ Replicas are placed as far from each other as possible
 - ▶ For example, if two regions are tied for least-used region, replicas are placed in two regions
- ▶ This process is called *ring-building*




<https://www.swiftstack.com/blog/2013/02/25/data-placement->

Consistency and Network Partitioning

2-node example


Akash 100
Agni takes 50
Akash 50



Akash 100

4. Bring servers down (no availability)
fly (inconsistent) OR

Practical example: Netflix




- ▶ Netflix: video on demand over the Internet
- ▶ Runs on Amazon cloud
- ▶ Consider the following scenario
 - ▶ User at TV updates list of favorites
 - ▶ Load balancer sends update to server 1
 - ▶ Set top box requests favorites list
 - ▶ Load balancer sends update to server 2
 - ▶ Is the returned result consistent?
Depends!
- ▶ Comparing NoSQL Availability Models by Adrian Cockcroft,
<http://perfcap.blogspot.com/2010/10/comparing-nosql-availability-models.html>

Implications of CAP Theorem

- ▶ Conventional view
 - ▶ Network partitioning will happen
 - ▶ Includes high latency
 - ▶ Cloud applications have to deal with either non-availability or inconsistency
- ▶ Contra view
 - ▶ CAP theorem focuses only on one cause of storage down time – human error, application error are equally important
 - ▶ Trade-offs should depend on system – mainframe or Windows
 - ▶ Partitions rare – not good to make general design decision based upon a rare case
 - ▶ Next-generation dbs are much faster – CAP theorem irrelevant

CAP theorem: diagram




Other replication architectures

- **Multi Leader Replication**
 - Use cases
 - Multi Data Center operation
 - Clients with offline operation – calendar apps on mobile phone
 - Collaborative Editing – Google Docs
- Leaderless Replication
 - No master node. Any node can potentially be master – Cassandra, Dynamo DB

Request Routing

- ▶ How does a client know which partition to get data from?
- ▶ Strategies
 - ▶ Client aware of partitioning strategy – so sends request directly
 - ▶ Client sends to routing tier – which routes to correct node
 - ▶ Client talks to any node which forwards the request.



Request Routing

- ▶ Use Zookeeper to coordinate knowledge of partition to node mapping

