

Hardware Virtualization



Virtualizable Architectures

Non-virtualizable Architectures

- ▶ x86 architecture is non-virtualizable
- ▶ What is the fundamental assumption that must hold if trap and emulate virtualization is to work?

Virtualizable Architectures

- ▶ The fundamental assumption for trap and emulate virtualization is
 - ▶ All sensitive instructions can be trapped
 - ▶ Sensitive instructions = instructions that *must* be emulated
- ▶ *Behavior sensitive*: result of instruction depends upon privilege level
 - ▶ If guest OS executes these instructions at lower privilege levels, result will be wrong
- ▶ *Control sensitive*: results that could change processor privilege level or change processor state
 - ▶ Like modify the processor registers like page tables.
- ▶ *Formal Requirements for Virtualizable Third-Generation Architectures by Gerald Popek and Robert Goldberg, Comm. ACM, 1974*

Exercise 1 (5 minutes)

- ▶ Consider an imaginary OS that executes the following imaginary instructions
 - ▶ MASK # turn off all interrupts before doing I/O
 - ▶ IO # do I/O
- ▶ MASK instruction works if executed in privileged mode
- ▶ If executed at non-privileged level, fails silently
 - ▶ Will fail if executed by imaginary OS in VM
- ▶ Would the imaginary OS work if run in a VM?
- ▶ Is MASK a (i) control sensitive (ii) behavior sensitive instruction
- ▶ Is the imaginary architecture virtualizable?

Non-virtualizable Aspects of x86 Architecture

Pops a doubleword (POPFD) from the top of the stack (if the current operand-size attribute is 32) and stores the value in the EFLAGS register, or pops a word from the top of the stack (if the operand-size attribute is 16) and stores it in the lower 16 bits of the EFLAGS register (that is, the FLAGS register). These instructions reverse the operation of the PUSHF/PUSHFD instructions.

as privileged as the IOPL. If a POPF/POPFD instruction is executed with insufficient privilege, an exception does not occur but privileged bits do not change.


- ▶ *popf* instruction
- ▶ Fails silently if guest executes it in non-privileged mode
- ▶ Cannot be handled by Trap-and-Emulate virtualization
 - ▶ There are about 17 instructions in x86 that are non-virtualizable

Binary Translation

Binary Translation


- ▶ Transparent virtualization technique
- ▶ Complements trap and emulate virtualization
 - ▶ Allows virtualization of x86 architecture
- ▶ Invented by VMWare
 - ▶ Mainframe architectures are virtualizable
- ▶ Conceptually simple
 - ▶ Replace *popf* instruction with trap to kernel and emulation
 - ▶ Replace simple instructions with equivalent instructions (details in paper)
 - ▶ Replacement is not done at run time; done before OS is run to create a modified OS image

What are Basic Blocks?



Source: <https://labs.vmware.com/download/45>, Scott Devine, VMWare

Binary Translation



Source: <https://labs.vmware.com/download/45>, Scott Devine, VMWare

Binary Translation Complexities


- ▶ Control flow changes
 - ▶ Translation can change addresses of instructions
 - ▶ Branches have to be re-translated
 - ▶ Keep track of branch addresses

```
isPrime:    mov    %eax, %edi ; %eax = %edi (a)
            mov    %esi, $2 ; i = 2
            cmp    %esi, %eax ; is i >= a?
            jge   prime      ; jump if yes
            ...
```

```
isPrime':   mov    %eax, %edi      ; IDENT
            mov    %esi, $2
            cmp    %esi, %eax
            jge   [takenAddr] ; JCC
            jnp   [fallthrAddr]
```

x86-32 workaround in VMWare Workstation

Picture from Sorav Bansal/Scott Devine



- ▶ **Idea:** Direct execution + Dynamic Binary translation
 - ▶ Direct execution of most user level instruction
 - ▶ Dynamic binary translation for most system code
 - ▶ Dynamically decide whether to use direct execution or


Workaround for x86-32: Paravirtualization

- ▶ Co-design of guest OS and hypervisor
 - ▶ Advantage: Simplicity, work around corner cases
 - ▶ Disadvantage: Need **modified** guest OS
 - ▶ Example: Xen hypervisor
- ▶ Trick to virtualize x86-32:
 - ▶ Block all 17 instructions that are sensitive but not privileged
 - ▶ Just modify the guest OS to #undef them
 - ▶ Instead provide **hypercalls** from guest OS to VMM
 - ▶ Hypercalls are like system calls but from guest OS to VMM


Hardware assisted virtualization: Intel VT-x/AMD-v for x86-64

- ▶ Solution idea:
 - ▶ Two new modes of operation: **root** and **non-root**
 - ▶ Each mode has complete set of execution rings (0-3)
 - ▶ New instructions to switch between modes
 - ▶ H/W state is duplicated for each operation mode
 - ▶ Hypervisor runs in root mode and VMs in non-root mode
 - ▶ When any “sensitive” instruction executed in non-root mode it either (1) executed by processor on duplicated state or (2) trap to hypervisor

Hardware assisted virtualization: Intel VT-x/AMD-v for x86-64



Using Intel VT-x/AMD-v (KVM way)




Current status of H/W assist

- ▶ Almost all VMMs make use of h/w assist today
- ▶ KVM is integrated part of Linux and makes heavy use of h/w assist
- ▶ Even Xen and VMWare workstation uses it


Memory Virtualization



Recap: Virtual Memory in un-virtualized system



Recap: Performing address translation in unvirtualized machine




Requirement for memory virtualization


- ▶ **No** direct access to physical memory from VM
- ▶ **Only** hypervisor should manage physical memory
 - ▶ Why?
- ▶ But, fool the guest OS to think that it is accessing physical memory

Virtualizing Virtual Memory


Managed by Virtual Machine
guest OS 1



Virtual Machine
2



Managed by
hypervisor




Host page table

Host (real)Physical memory

Virtualizing Virtual Memory

Based on "Efficient Memory Management" Gandhi et.
al., MICRO'14



Two levels of address translation on each memory access by an application running inside a VM

Implementing two level address translation

1

- ▶ Shadow page table
 - ▶ Software only technique

2

- ▶ Nested/Extended page table
 - ▶ Hardware support


Shadow page table

- ▶ **Idea:** Let hypervisor “create” a shadow page table that maps guest VA to host PA directly
 - ▶ Made by combining guest page table w/ system page table
 - ▶ Hypervisor makes the *cr3* point to the shadow page table


Shadow Page Tables

- ▶ VMM maintains shadow page tables that map guest-virtual pages directly to machine pages.
- ▶ Guest modifications to V->P tables synced to VMM V->M shadow page tables.
 - ▶ Guest OS page tables marked as read-only.
 - ▶ Modifications of page tables by guest OS -> trapped to VMM.
 - ▶ Shadow page tables synced to the guest OS tables

Set CR3 by guest OS (1)



Set CR3 by guest OS (2)



Drawbacks: Shadow Page Tables


- ▶ Maintaining consistency between guest page tables and shadow page tables leads to an overhead: VMM traps
- ▶ Loss of performance due to TLB flush on every “world-switch”.
- ▶ Memory overhead due to shadow copying of guest page tables.

Review

- ▶ True or false: x86 architecture is virtualizable
- ▶ What is binary translation?
- ▶ False
- ▶ Replacement of sensitive instructions by traps

I/O Virtualization

I/O transfers on Baremetal




Physical address I/O: Challenges


- ▶ Can't swap the page out
- ▶ External device can write into any memory location – Security concerns

- ▶ How will you solve this problem?


I/O transfers on Baremetal



VMM Programs IOMMU




VM programs IOMMU




Options For I/O Virtualization


Monolithic Model



Service VM Model



Pass-through Model



- Pro: Higher Performance
- Pro: I/O Device Sharing
- Pro: VM Migration
- Con: Larger Hypervisor

VMWare
ESX

- Pro: High Security
- Pro: I/O Device Sharing
- Pro: VM Migration
- Con: Lower Performance

Xen

Shared Device can be file, disk or disk partition

- Pro: Highest Performance
- Pro: Smaller Hypervisor
- Pro: Device assisted sharing
- Con: Migration Challenges